### **Operating Systems**

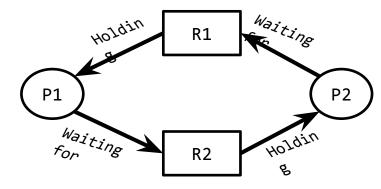
# **Deadlock**

### Deadlock

A set of processes is in a deadlocked state if every process in the set is **waiting for an event** that can be caused only by another process in the set.

#### Example:

- Two process : P1, P2
- Two resource: R1, R2
- P1 is holding R1 and waiting for R2
- P2 is holding R2 and waiting for R1



# System Model

- A system contains finite number of resources
  - These resources are distributed among competing processes
  - Examples: CPU cycle, File, I/O devices ( printer, DVD drive) etc.
- Each resource type may have identical number of instances
- A process may request as many resource as it requires.
  - But, it should not exceed the total number of available resources of the system
- A process may utilize a resource only in the following sequence
  - **1. Request**: requests the resource
  - **2. Use**: operate on the resource
  - **3. Release**: releases the resource

# Necessary Conditions for Deadlock

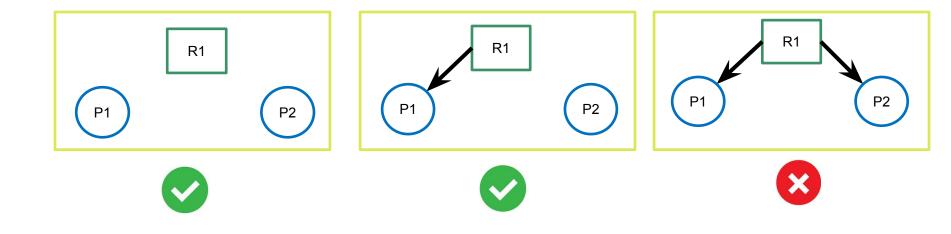
A deadlock situation can arise if the following four conditions hold simultaneously in a system -

- 1.Mutual exclusion
- 2.Hold and wait
- 3.No preemption
- 4.Circular wait

If one of them is not present in a system, no deadlock will arise

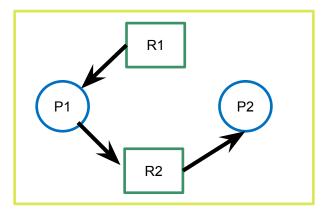
### Mutual Exclusion

At least one resource must be held in a nonsharable mode; that is, only one process at a time can use the resource



### Hold and Wait

A process must be holding at least one resource, and waiting to acquire additional resources that are currently being held by other processes.



# No Preemption

a resource can be released only voluntarily by the process holding it, after that process has completed its task.





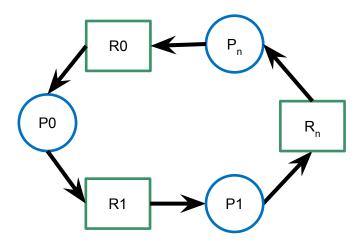
### Circular Wait

A set  $\{P_0, P_1, ..., P_n\}$  of waiting processes must exist such that

P0 is waiting for a resource held by P<sub>1</sub>,

 $P_1$  is waiting for a resource held by  $P_2$ , ...,  $P_{n-1}$  is waiting for a resource held by  $P_n$ , and

 $P_n$  is waiting for a resource held by  $P_0$ 



# Resource Allocation Graph

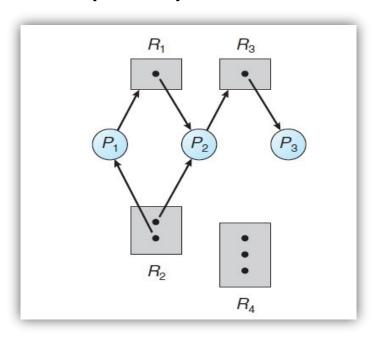
# Resource Allocation Graph

A set of vertices V and a set of edges E.

- V is partitioned into two types:
  - $P = \{P_1, P_2, ..., P_n\}$ , the set consisting of all the processes in the system.
  - $\circ$  R = {R<sub>1</sub>, R<sub>2</sub>, ..., R<sub>m</sub>}, the set consisting of all resource types in the system.
- Request edge: directed edge from a process to a resource
  - $\circ$   $P_1 \rightarrow R_2$
- Assignment edge: directed edge from a resource to a process
  - $\circ \quad R_1 \to P_2$

Process	$P_i$
Resource Type with 4 instances	
$P_{i} \text{ requests instance} \\ \text{ of } R_{j}$	$P_i$
$P_i$ is holding an instance of $R_j$	

# Example of Resource Allocation Graph

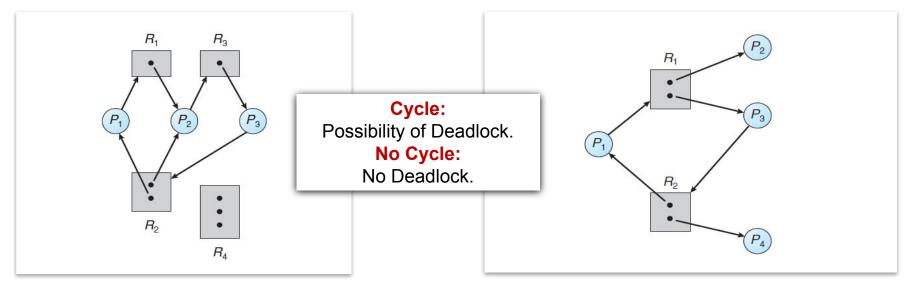


$$P = \{P1, P2, P3\}$$

$$R = \{R1, R2, R3, R4\}$$

$$E = \{P1 \rightarrow R1, P2 \rightarrow R3, R1 \rightarrow P2, R2 \rightarrow P2, R2 \rightarrow P1, R3 \rightarrow P3\}$$

### Resource Allocation Graph (Deadlock)



Cycle 1:  $P1 \rightarrow R1 \rightarrow P2 \rightarrow R3 \rightarrow P3 \rightarrow R2 \rightarrow P1$ 

Cycle:  $P1 \rightarrow R1 \rightarrow P3 \rightarrow R2 \rightarrow P1$ 

Cycle 2:  $P2 \rightarrow R3 \rightarrow P3 \rightarrow R2 \rightarrow P2$ 

No Deadlock!

**Deadlock!** 

# Methods for Handling Deadlocks

#### Prevention:

• Ensure that the system will never enter a deadlock state

#### Avoidance:

request for any resource will be granted if the resulting state of the system doesn't cause deadlock

#### Detection and recovery:

• Allow the system to enter a deadlock state and then recover

Operating Systems

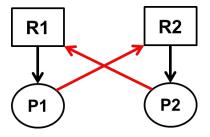
# Deadlock Prevention

### Deadlock Prevention

Ensure that at least one of the necessary conditions for deadlock cannot hold in the system.

#### **Mutual Exclusion:**

- Mutual exclusion condition must hold
- At least one resource must be non-sharable
- Sharable resources do not require mutually exclusive access (example: Read only files)



#### **Hold and Wait:**

Guarantee that when a process requests a resource, it does not hold any other resources.

- <u>Protocol 1:</u> Request and allocate all its resources before execution.
- <u>Protocol 2:</u> Allow process to request resources only when the process has none.

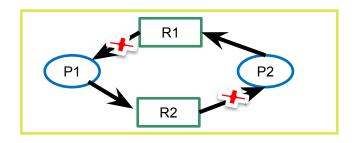
#### Disadvantages:

- Low utilization of resources
- Starvation

### Deadlock Prevention

#### **No Preemption:**

- If a process is holding some resources and requesting for another resource, which is held by another process, then this process needs to wait, and the resources held by it will be preempted or released.
- Preempted resources are added to the list of resources for which the process is waiting
- Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting.



#### Circular Wait:

- Impose a total ordering of all resource types
- Each process requests resources in an increasing order of enumeration

 $R = \{R_1, R_2, ..., R_m\}$  and we assign unique number to each of the resource type.

For example,  $F(R_1) = 1$ ,  $F(R_2) = 3$ ,  $F(R_3) = 5$ .

- A process can initially request any number of instances of a resource type,  $R_i$
- After that, the process can request instances of resource type  $R_i$  if and only if  $F(R_i) > F(R_i)$ .
- A process requesting an instance of resource type  $R_j$  must have released any resources  $R_i$  such that  $F(R_i) \ge F(R_i)$

Operating Systems

# Deadlock Avoidance

### **Deadlock Avoidance**

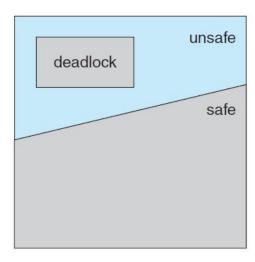
- System has a priori information, such as, maximum number of resources of each type that every process need
- Based on these information construct an algorithm that ensures that the system will never enter a deadlocked state
- Dynamically examines the resource-allocation state to ensure that a circular-wait condition can never exist
- Resource-allocation state: number of available and allocated resources, and the maximum demands of the processes.

### Safe State

- System is in safe state if can allocate resources to each process (up to its maximum) in some order and still avoid a deadlock.
- Safe state = there exists a safe sequence
- Sequence  $\langle P_1, P_2, ..., P_n \rangle$  is safe if, for each  $P_i$ , the resources that  $P_i$  can still request, can be satisfied by (currently available resources + resources held by all the  $P_i$ , with j < i)
  - If  $P_i$  resource needs are not immediately available, then  $P_i$  can wait until all  $P_i$  have finished.
  - When  $P_j$  is finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and terminate.
  - When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on.

### Safe State, Unsafe State & Deadlock

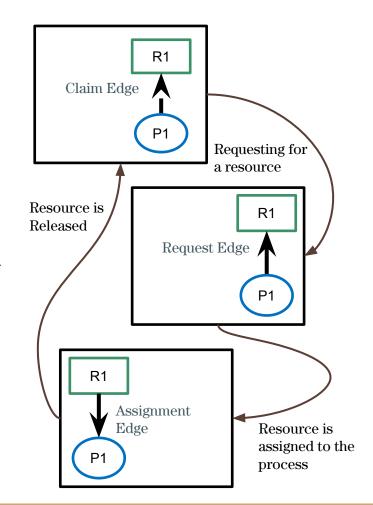
- Safe State => No Deadlock!
- Unsafe State => Possibility of Deadlock!
- To avoid deadlock, always ensure that system is in safe state.



### Resource-Allocation-Graph Algorithm

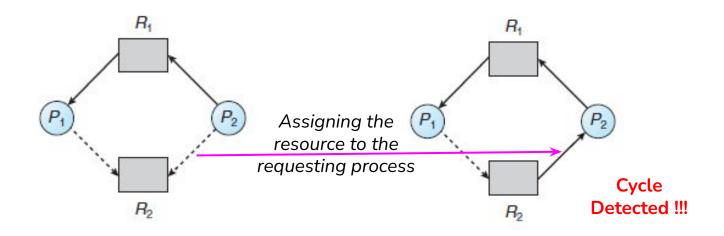
(Deadlock Avoidance Algorithm)

- Used when the system has only one instance of each resource type.
- Request Edge and Assignment Edges exist here similar to Resource allocation graph
- Claim Edge: claim edge  $P_i \rightarrow R_j$  indicates that process  $P_i$  may request resource  $R_i$  at some time in the future
- When process  $P_i$  requests resource  $R_j$ , the claim edge  $P_i \rightarrow R_j$  is converted to a request edge.
- When a resource  $R_j$  is released by  $P_i$ , the assignment edge  $R_j \rightarrow P_i$  is reconverted to a claim edge  $P_i \rightarrow R_j$
- Resources must be claimed a priori in the system. i.e. before process  $P_i$  starts executing, all its claim edges must appear in the graph



## Resource-Allocation-Graph Algorithm

(Deadlock Avoidance Algorithm)



Resource-allocation graph for deadlock avoidance

An unsafe state in a resource-allocation graph

### **Operating Systems**

# Bankers Algorithm

### Banker's Algorithm (Deadlock Avoidance Algorithm)

- Multiple instances of each resource type.
- Each process must a priori claim maximum use.
- When a process requests a resource it may have to wait.
- When a process gets all its resources it must return them in a finite amount of time.

# Data Structures for the Banker's Algorithm

Need the following data structures, where  $\mathbf{n}$  is the number of processes in the system and  $\mathbf{m}$  is the number of resource types.

- Available: Array of length **m**, indicates the number of available resources of each type.
- Max: **n** by **m** matrix, defines maximum demand of each process for each type.
- □ Allocation: **n** by **m** matrix, defines number of assigned resources to each process for each type.
- Need: **n** by **m** matrix, defines number of remaining resources of each type needed by each process.

```
Need [i,j] = Max[i,j] - Allocation [i,j]
```

# Banker's Safety Algorithm

Algorithm for finding out whether or not a system is in a safe state

Let Work and Finish be vectors of length m and n, respectively. Initialize Work = Available and Finish[i] = false for i = 0, 1, ..., n - 1.

- Find an index i such that both of the following meets, Finish[i] == false Need₁ ≤ Work if no such i exists, go to step 4.
- **3.** Work =Work + Allocation<sub>i</sub> Finish[i] = true
  Go to step 2.
- 4. If **Finish**[i] == **true** for all *i*, then the system is in a safe state.

### Banker's Algorithm Example

- 5 processes P0 through P4;
- 3 resource types A (10 instances), B (5 instances), and C (7 instances)
- at time T0, the following snapshot of the system:

	Allocation	Max	Available		Need
	ABC	ABC	ABC		ABC
$P_0$	010	753	332	$P_0$	743
$P_1$	200	322	40 10 00 00 00 00 00 00 00 00 00 00 00 00	$P_1$	122
$P_2$	302	902		$P_2$	600
$P_3$	211	222		$P_3$	011
$P_4$	002	433		$P_4$	431

inish:			
work:			

The system is in a safe state since the sequence  $< P_1, P_3, P_4, P_2, P_0 >$  satisfies safety criteria

# Practice

	Allocation				Max			Available		
	A	В	C	A	В	C	A	В	C	
$P_0$	1	0	1	2	1	1	2	1	1	
$\mathbf{P}_1$	2	1	2	5	4	4				
P <sub>2</sub>	3	0	0	3	1	1				
P <sub>3</sub>	1	0	1	1	1	1				

# Resource-Request Algorithm for Process P<sub>i</sub>

- 1. If  $Request_i \leftarrow Need_i$  go to step 2.
  - Otherwise, raise error condition, since process has exceeded its maximum claim.
- 2. If **Request**<sub>i</sub> <= **Available**<sub>i</sub>, go to step 3.
  - Otherwise  $P_i$  must wait, since resources are not available.
- 3. Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:

```
Available = Available - Request;
Allocation; = Allocation; + Request;
Need; = Need; - Request;
```

- 4. Check Bankers Safety Algorithm, if this new state is safe.
  - $\rightarrow$  If safe the resources are allocated to  $P_{t}$
  - $\rightarrow$  If unsafe  $P_i$  must wait, and the old resource-allocation state is restored

# What if P<sub>1</sub> Request (1,0,2) more?

- To decide whether this request can be immediately granted,
  - we first check that Request<sub>1</sub>  $\leq$  Need<sub>1</sub> that is,  $(1,0,2) \leq (1,2,2)$ ? => TRUE
  - Then we check that Request<sub>1</sub>  $\leq$  Available that is,  $(1,0,2) \leq (3,3,2)$ ? => TRUE
  - If both is true,

Available = 
$$(3,3,2) - (1,0,2) = (2,3,0)$$
  
Allocation<sub>1</sub> =  $(2,0,0) + (1,0,2) = (3,0,2)$   
Need<sub>1</sub> =  $(1,2,2) - (1,0,2) = (0,2,0)$ 

• So, we arrive at the following new state

Is the system in safe state? Yes. Because we find a safe sequence:  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ . Grant the request of  $P_1$ .

What about request of (3 3 0) resources for P4? What about request of (0 2 0) resources for P0?

	Allocation	Need	Available
	ABC	ABC	ABC
$P_0$	010	743	230
$P_1$	302	020	
$P_2$	302	600	
$P_3$	211	011	
$P_4$	002	431	

